



Department of Mathematics and Computer Science

Enhancing Face Detection and Clustering for Image Analysis through an Interactive Visualization: A Broader Scope

Bachelor End Project Report

Thijmen Broeren - 1839527

supervisors:

Huub van de Wetering

1-7-2025

Abstract

Face detection and clustering have always been challenging tasks since their practical use cases can differ by a lot and these models lack the ability to generalize. The problem addressed consists of assessing how more reliable face detection models and clustering algorithms can be created, and how an interactive visualization can provide more insight into the model's correctness. Building on prior work, face detection and clustering models are improved and evaluated. For face clustering, various models of the YOLO architecture are used, a self-trained model is created using the newest YOLO architecture, and an InsightFace model is used. From these models, the self-trained model performs best. For face clustering, the DBSCAN algorithm and Chinese Whispers algorithm are tested with different parameters, and seven different face embedding models are tested of which the combination of the DBSCAN algorithm and the face_recognition model performed best. These models are then used in a visual analytics tool. This tool allows a user to interact with these models to create an understanding of their functioning and use them for analysis. The main purpose of this analytical tool is to allow a user to bring their own set of images to find common co-occurring person within that dataset, which is, among others, useful for political image analysis. Moreover, additional data, like categories, are allowed to provide more context to the user if needed. The code for all experiments and the interactive visualization can be found on the GitHub¹.

¹<https://github.com/stiverthijmen163/BEP/tree/main>

Contents

1	Introduction	3
2	Related Work	5
3	Datasets	7
3.1	Face Detection	7
3.2	Face Clustering	9
3.3	Visualization	9
4	Methodology	11
4.1	Face Detection	12
4.2	Face Clustering	15
5	Interactive Visualization	18
5.1	User Tasks	18
5.2	Design Sketch	19
5.3	Design Choices - Visuals	21
6	Results	24
6.1	Face Detection Model Selection	24
6.2	Face Clustering	25
6.3	Interactive Visualization	26
7	Conclusion and Discussion	32
7.1	Face Detection	32
7.2	Face Clustering	32
7.3	Interactive Visualization	32
A	Appendix	37
A.1	Pseudocode	37
A.2	Dimensionality Reduced Clustering	38
A.3	t-SNE Reduced Embeddings with Different Models	38
A.4	Testing Parameters for Clustering with DBSCAN	42
A.5	Testing Parameters for Clustering with Chinese Whispers	48

1 Introduction

Understanding how people interpret political news is influenced by multiple factors, including the way the news article is presented. When reading through political news articles, it can be interpreted differently among people. This is because of textual and visual modalities within the news. By analyzing textual and visual modalities, we can for example research whether there are certain politicians who are overrepresented compared to others within the news pages (Ruyters, 2025). Ruyters has proposed a method for this by using face detection models to detect and isolate the faces from visual modalities, and by clustering the faces (one cluster/class per unique face) as a form of face recognition. Not only can any overrepresented classes be found this way, but also any classes commonly co-occurring in visual modalities. Using the proposed solution, one could achieve more insights into the impact of textual and visual modalities on the interpretation of news content. This thesis builds on that idea by focusing on visual modalities to identify commonly co-occurring persons in news articles.

To understand the impact of visual modalities on the interpretation of news content it is, among others, important to have an insight as to whom are represented in representative images at the start of a news article (picture lead). While these detection models already exist to some extent, they tend to be unreliable and not robust. Therefore, this study aims to provide more reliable models and more insight into how these models work and what their strengths and limitations are. These improvements aim to ensure a more accurate representation of individuals in political news images.

Face detection and recognition have become more and more relevant over the years, not only in the context of politicians in the news. Therefore a study such as this can apply to multiple applications, which is why the end product, an interactive visualization, should be easily applicable to multiple applications. For instance, to find co-occurring persons in sports images. However, the scope is also broadened to future-proof the end product in the sense of being able to add data as a user with additional information such as topics or the most common words in a news article. This allows a user to perform deeper analysis, not only looking at co-occurring persons, but also keeping track of some additional information for those co-occurring persons.

Obviously, when a model is unreliable, it can lead to wrong conclusions. This can have severe consequences if these conclusions are used for the representativeness of politicians in news picture leads since this tends to prove that certain news pages may be biased toward certain politicians. To prevent this from happening robustness and reliability are two important requirements for algorithms since they allow a user to be confident that the result is an accurate representation of reality. Apart from creating better models, a way of improving this is by providing insight to the user as to how well a certain detection/clustering algorithm performs and allowing the user to test them in different circumstances such as noisy data or images taken at night. This kind of transparency can be achieved through an interactive visualization like the one proposed in this thesis.

In summary, this study focuses on improving face detection and clustering models compared to the methods introduced by Ruyters (2025) by creating different models and

using more recent techniques. In addition, an interactive visualization is created to provide more insight into the model's performance and therefore the reliability of the model. An extra, but important feature, is that this study attempts to create a visualization such that it applies to different fields of study.

2 Related Work

The aim of this thesis lies in the improvement of the existing face detection and clustering performance as well as within the development of a visual analytics tool that provides more insight into the model’s reliability. It therefore combines a number of research areas within computational sciences, namely: 1) object detection, 2) face clustering, and 3) visual analytics. This section discusses these fields of research and their application to this thesis in more detail.

Object Detection: According to Ali and Zhang (2024), one way of doing object detection is by using the You Only Look Once (YOLO) architecture. There are many versions of YOLO you could use for this all performing with $mAP > 0.90$. mAP is the mean of the average precision across all classes in a model, for more information see Solawetz (2020). Also, in Ruyters (2025) two different methods are used which do not work very well apart from each other but have increased accuracy when combined, namely `haar_cascade` and `face_recognition`. These will be used as a base model. Also, the technique of combining different models to get more detections in the database is used here. Another technique used in this thesis is InsightFace which is known for its models in face detection and recognition with high accuracy (Guo et al., 2021).

Face Clustering/Classification: When working with unlabelled data it is important to cluster the objects found by the detection models. To do this one could use K-nn and SVM as a pre-trained model and run it on the unlabelled data, but this leads to some misclassifications (Ruyters, 2025). Moreover, when clustering unlabelled data and there does not exist such a recognition model yet, K-nn or SVM would be no option. To solve this problem, there are two important factors: the embeddings used for clustering and the face clustering algorithm (Petukhova et al., 2025).

For clustering, the Chinese Whispers algorithm has many applications (Biemann, 2006; Shiells and Pham, 2010; Al Ali et al., 2024). However, according to Singh and Hang (2021) and Lin et al. (2021), it can also be used for face classification and it is considered a relatively quick algorithm. One problem that arises though is that this algorithm performs better with larger datasets. The scikit-learn library also has many easy-to-implement clustering algorithms, one of which is the DBSCAN algorithm (Pedregosa et al., 2011). DBSCAN is a density-based clustering algorithm that works well for all kinds of shapes as long as the data points are close to each other. There are many ways to embed faces, one of which is FaceNet using face similarity to embed faces (Schroff et al., 2015). DeepFace introduces a way to easily implement FaceNet and some other face embedding models along with it such as VGG-Face and OpenFace (Serengil and Ozpinar, 2021). One problem that arises though is that DeepFace does not allow for inputting the locations of a face, therefore the effectiveness of a good face detection model may be reduced. One model that does allow this, is the `face_recognition` model (Geitgey, 2020).

Visual Analytics: The end product of this thesis will be a tool for visual analytics, i.e. an interactive visualization facilitating analytical reasoning (Thomas and Cook, 2005). It uses a front-end to allow a user to change parameters, interact with visuals, and analyse the data to their liking. The back-end is used to adapt to the changes by the user by running some analysis on some data (Preim et al., 2023). For this thesis, the visual analytics tool will be used to interpret and influence the results coming from both the face detection model and the face clustering algorithm.

Moreover, visualizations allow users to actively engage with the models, resulting in a better understanding of the models (Grissom et al., 2003; Naps, 2005). According to Saraiya et al. (2004) and Munzner (2014), there are specific features that are most effective for such visualizations. Depending on the context, this means for example that a bar chart may be a better visualization instead of a pie chart and that certain colours like green and yellow should be avoided within the same figure.

3 Datasets

To train, validate, and evaluate models and algorithms worked with in this thesis, various datasets are needed. This section aims to explain the different datasets used for each of the different aspects within this thesis; 1) face detection, 2) face clustering, and 3) a visual analytics tool.

3.1 Face Detection

For training and validating face detection models, a subset of two datasets is used: the Face Detection Dataset² and the WIDER FACE dataset³. The Face Detection Dataset consists of 13389 train and 3347 validation images. The images are high-quality images with one or more faces in each of them and in a large variety of scenes. The WIDER FACE dataset consists also of high-quality images, typically images contain a larger amount of faces with a greater distance from the camera divided into 61 different types of scenes. It has a total of 12880 train and 3226 validation images. Both datasets come with the bounding boxes for all faces in their respective images. Some example images are shown in Figure 1. Using both these datasets the following two datasets are created:

FDD-WF small subset: This dataset is used for training face detection models in a quick matter. It consists of 500 random training images from both datasets and 100 random validation images from both datasets.

FDD-WF subset: This dataset is a larger version of the 'FDD-WF small subset' dataset such that it can be used for training a final version of the face detection model. It contains 2500 random training images from both earlier mentioned datasets and 500 random validation images from both those datasets.

²<https://www.kaggle.com/datasets/fareselmenshawii/face-detection-dataset>

³https://huggingface.co/datasets/CUHK-CSE/wider_face



Figure 1: Two examples of the Face Detection Dataset (left) and two of the WIDER FACE dataset (right), including the bounding boxes.

FD-test: For the test data, a subset of the images scraped from the internet by Ruyters has been used. Only the images containing politicians are used to significantly decrease the size of the test data, this is done because these images must be labelled by hand, which is explained in section 4.1.2. Therefore, this dataset contains 290 images used in news headers. Some examples are shown in Figure 2.



Figure 2: Two examples of the original images used for the FD-test dataset, original images from <https://www.nu.nl/vragenoproep/6289454> & <https://nos.nl/artikel/2491240>.

3.2 Face Clustering

CFI: The Celebrity Face Image Dataset⁴ contains high-quality photos of 17 distinct celebrities. This is overall a balanced dataset, all classes except one have 100 images each, the other class has 200 images.

Embedding on the Wall: As a second dataset for face clustering, a more realistic and less ideal dataset is used. Ruyters created a dataset⁵ with all faces cut out from the news images, this includes a class called 'unknown' which will not be used, since this adds too much noise to the dataset to establish a good result when clustering. This dataset is unbalanced and contains images both of low and high quality and resolution. The number of occurrences per class ranges from 1 (Edson Elf) to 66 (Pieter Omtzigt).



Figure 3: Four examples of the Celebrity Face Image Dataset (left) and four of the Embedding on the Wall dataset (right).

3.3 Visualization

To test and showcase the functionalities of the interactive visualization, an extra dataset is necessary. Note that the tool can have two types of input data, a set of images or a CSV-file. This CSV-file must contain of images encoded using base64 to reduce the size, and can contain some additional data. This additional data can be some lists of attributes, or for example a URL. This section explains one such CSV-file to use for the tool.

⁴<https://www.kaggle.com/datasets/vishesh1412/celebrity-face-image-dataset>

⁵https://github.com/wieswies/embedding_on_the_wall/tree/main/d_nos-nu-news-images-classification/datasets/images/svc_classification_corrected

Tool-Test: This dataset, as shown in Figure 4, uses the same images as the FD-test dataset, excluding the labels for all faces. In addition, it contains two additional information columns and a URL column. The *category* column contains the categories for each news article and the *Issues_paragraphs* column contains the issues handled for each paragraph in that news article. Both could be interesting for a political image analyst to acquire knowledge as to which politician is associated with which categories and issues.

id	url	category	Issues_paragraphs	img
2488763	https://www.nos.nl/...	["Binnenland", "Buitenland"]	[]	/9j/4AAQSkZ...
2488778	https://www.nos.nl/...	["Op weg naar TK2023", "Politiek"]	["migratie", "arbeidsmigranten", "bestaanszekerheid", "bestaanszekerheid"]	/9j/4AAQSkZ...
2488800	https://www.nos.nl/...	["Op weg naar TK2023", "Politiek"]	[]	/9j/4AAQSkZ...
2488814	https://www.nos.nl/...	["Binnenland"]	[]	/9j/4AAQSkZ...
2488823	https://www.nos.nl/...	["Op weg naar TK2023", "Politiek"]	[]	/9j/4AAQSkZ...
2488855	https://www.nos.nl/...	["Op weg naar TK2023", "Politiek"]	["bestaanszekerheid", "woning", "bestaanszekerheid"]	/9j/4AAQSkZ...
2488891	https://www.nos.nl/...	["Op weg naar TK2023", "Politiek"]	["woning", "woning", "migratie"]	/9j/4AAQSkZ...
2488894	https://www.nos.nl/...	["Binnenland", "Buitenland"]	[]	/9j/4AAQSkZ...
2488934	https://www.nos.nl/...	["Op weg naar TK2023", "Politiek"]	[]	/9j/4AAQSkZ...
2488941	https://www.nos.nl/...	["Op weg naar TK2023", "Politiek"]	["Immigratie"]	/9j/4AAQSkZ...

Figure 4: The first ten rows from the Tool_Test dataset, displaying what the user input would look like when adding additional information and URL's.

4 Methodology

The main result of this thesis consists of an interactive visualization in which users will be able to understand how well the models and algorithms perform and are able to influence the results of such models, as explained in section 5. To come to this visualization, a few steps are needed to create such models and algorithms, this is visualized in Figure 5.

The proposed method consists of three steps which are applied in sequence to create the visualization. The first step is the creation of a face detection model with one or more images as input. The output of such a model consists of images of all faces within the input images. Here different models are trained, tested, and compared to the model used in Ruyters (2025). The best of these models is chosen as explained in section 4.1 and is used in the interactive visualization.

The second step is the creation of a face clustering algorithm which should be able to cluster all faces in a given set of faces such that every unique face has its own cluster. Within the final algorithm, the input images for this algorithm are the output images of the face detection model. Also, multiple cluster algorithms will be implemented of which the best algorithm as explained in section 4.2 is chosen to be used in the interactive visualization.

Finally, a face can be inserted into the visualization of which the user would want the results as explained in section 5. The model should then be able to apply face recognition and find the correct cluster that matches the input face if such a cluster exists. To do so, a simple face recognition will be fitted on the output of the clustering algorithm as explained in the paper of Ruyters, and then be used to predict the cluster for the input image.




Figure 5: Flowchart displaying the approach of the main steps of the methodology.

4.1 Face Detection

Even though face detection, i.e. given an input image returning bounding boxes defining the coordinates of faces in that image, is considered a solved problem within computational sciences, many different approaches co-exist in practice. This thesis tests and assesses five; 1) a base model consisting of a combined pipeline with face_recognition and HAAR cascades, 2) InsightFace, 3) YOLOv6n-face, 4) YOLOv611l-face, and 5) YOLOv12-face. This section briefly explains the differences between the mentioned models.

1. **Base model:** This is the face detection model as used in the paper of Ruyters (2025), combining two techniques:
 - (a) **Face_recognition**⁶ is a face recognition model based on dlib's face recognition model⁷. According to their package page, it has an accuracy of 99.38% on the Labelled Faces in the Wild dataset, but their model can also be used to detect faces instead of face recognition.
 - (b) **Haar-cascade**⁸ is a model with which multiple objects can be detected like human bodies, faces, and eyes. For face detection, haar-cascade offers a few models that can be used. In this thesis, the same model as used by Ruyters will be used.

Ruyters combines these face detection models to increase the number of detected faces and removes the overlapping faces between the two models.

2. **InsightFace:** This is an open-source deep face analysis library, offering multiple face-detection models. The buffalo_l model is one of those that uses the SCRFD-10GF architecture, designed to optimize both speed and accuracy (Guo et al., 2021). It reaches an accuracy of 95.16%.
3. **YOLOv6n-face & YOLOv611l-face:** These models use the YOLO architecture as described by Jocher et al. (2023), adapted for face detection purposes. On the models' GitHub page⁹, all different model versions trained for face detection can be found. For the purpose of this thesis, the oldest smallest model prioritizing speed is used (YOLOv6n-face) and the largest most recent model prioritizing accuracy is used (YOLOv11l-face).
4. **YOLOv12-face:** Recently, a new version of YOLO came out which has not been implemented as a face detection model yet¹⁰. Therefore, an additional YOLOv12-face model is implemented using the new YOLOv12 architecture similar to the two previously named models.

⁶https://github.com/ageitgey/face_recognition

⁷<https://dlib.net/>

⁸<https://github.com/opencv/opencv/tree/master/data>

⁹<https://github.com/akanametov/yolo-face>

¹⁰Checked at 01-03-2025

4.1.1 YOLOv12 model evaluation

The different versions of the YOLOv12 architecture (-n, -s, -m, -l, and -x) differ significantly, for which this thesis tests them individually. Since training a model at full scale requires a significant amount of time, a smaller dataset (FDD-WF small subset) was used to train and assess the model's performance using the following metrics:

- **mAP50:** The mAP calculated at an Intersection over Union (IoU) threshold of 0.50. The calculations of the IoU will be explained further on in this section.
- **mAP50-95:** The average of mAP calculated at different IoU thresholds between 0.50 and 0.95.
- **precision:** The proportion of predicted faces that are actually correct.

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}} \quad (1)$$

- **recall:** The proportion of all actual faces that are correctly classified as faces.

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad (2)$$

- **runtime:** The time needed for training the model.

Using these evaluation metrics, the YOLOv12 model can be chosen that learns the fastest and most accurately. Using this chosen model, a new model is trained on that architecture with a larger dataset, which is the FDD-WF subset dataset. To train this model, similar parameters are passed through as how the YOLOv6n-face - YOLOv11l-face models are trained.

4.1.2 Testing

For all models mentioned above, the same test data is used. Since the intended use for these models is the identification of faces in picture leads in newspapers considering politics, a representative dataset should be used. Therefore, the data scraped from the internet by Ruyters has been used and reduced to a set of 290 images such that every image contains some politician, referred to as the FD-test dataset. These images have been annotated by hand as to what you should expect a face detection model to recognize. An example of this can be seen in Figure 6, the entire test set can be found at the Roboflow project (BEP, 2025).



Figure 6: Example image of annotated test data for face detection. Original image from: <https://www.nos.nl/artikel/2489302>

To choose the best model among those tested, various evaluation metrics are used to check how the models perform on a representative dataset:

- **IoU (entire image):** The Intersection over Union (IoU) calculated over the entire image as shown in Equation 3. Shows how well the model's prediction covers the actual faces.
- **Accuracy (IoU ≥ 0.5):** Calculates the accuracy with an IoU threshold set at 0.5. The threshold of 0.5 is generally used as a base since it means that a face is detected but that the bounding box doesn't fit the actual bounding box well. The pseudocode for this calculation can be found in Algorithm 1.
- **Accuracy (IoU ≥ 0.7):** Calculates the accuracy with an IoU threshold set at 0.7. The threshold of 0.7 is used because it allows to check for predicted bounding boxes that fit the actual bounding box better.
- **Latency:** The speed of the model per image

$$\text{IoU} = \frac{\text{\#pixels where both a face is predicted and it actually is (overlap)}}{\text{\#pixels where either a face is predicted or one actually is (union)}} \quad (3)$$

4.2 Face Clustering

Considering the problem of face clustering, there are many different approaches to solving this problem. Limiting the options, the clustering model needed for this thesis must meet one requirement: being able to cluster an unknown set of faces, i.e. images resulting from all bounding boxes found by a face detection model, such that each unique face has its own cluster and all non-unique faces belong to the same cluster. The face clustering model considered in this thesis, consists of two phases:

1. **Embedding faces:** Vector representations are required for each face, capturing facial features using a numerical format such that similar faces are represented by similar vectors.
2. **Clustering Embeddings:** Using the calculated embeddings, a clustering algorithm in accordance with the specified requirement attempts to create unique classes for similar face embeddings.

This section describes and explains the approach to the mentioned phases.

4.2.1 Face Embeddings

Calculating embeddings can have many approaches differing in quality and vector size, therefore significantly impacting the resulting clustering by a clustering algorithm. This thesis tests and assesses seven embedding models through two packages, including FaceNet as used by Ruyters:

1. **face_recognition:** Applied for face detection, this package can additionally be utilized for calculating face embeddings. Furthermore, it allows a user to specify the bounding box of a face.
2. **DeepFace:** This package includes a feature to calculate face embeddings with six different models, namely: VGG-Face, FaceNet, OpenFace, DeepFace, Dlib, ArcFace (Serengil and Ozpinar, 2024, 2020). However, the user can't specify the location of a face, which may limit the quality of these embeddings.

4.2.2 Face Clustering

As explained, the clustering problem has many solutions. Ruyters solved this by creating a pipeline trained on some labelled face embeddings, which does not meet the specified requirement for the face clustering model. This thesis implements two clustering algorithms that are optimized, tested, and assessed; 1) Chinese Whispers, and 2) Density-Based Spatial Clustering of Applications with Noise (DBSCAN). Below, the functioning of both algorithms is explained.

1. **Chinese Whispers:** This algorithm creates a graph such that each embedding has its own vertex. After which, for each pair of vertices, the similarity is calculated

as shown in Equation 4, determining the weight for the edge between that pair of vertices. The algorithm then iteratively updates the clustering by appending the neighbours of a cluster to that cluster, considering only those edges above a chosen threshold.

Let E denote the set of face embeddings, and let D denote the set of Euclidean distances between all distinct pairs of embeddings in E :

$$D = \{\|\text{embedding}[u] - \text{embedding}[v]\| \mid u, v \in E, u \neq v\}$$

The Similarity between two faces i and j is then defined as follows:

$$\text{Similarity}(i, j) = 1 - \frac{\|\text{embedding}[i] - \text{embedding}[j]\|}{\max(D)} \quad (4)$$

2. **DBSCAN:** This algorithm introduces a new type of point: the core point. A point is a core point only if it meets the requirements set by the following parameters:

- **epsilon:** This parameter specifies the maximum distance between two points to be considered to be in the neighbourhood of each other.
- **min_samples:** Considering only those points within distance epsilon, a point is a core point only if it has at least min_samples neighbouring points.

To find the clusters, the DBSCAN algorithm puts all core points within distance epsilon of each other within the same cluster. For all remaining points, the point belongs to the cluster where a core point is within distance epsilon. If such a point does not exist, it belongs to the *unknown* cluster. To comply with the requirement of the clustering model as defined by this thesis, min_samples is set to one. This implies that every point is a core point and therefore secures that no point is added to the *unknown* cluster.

4.2.3 Testing

The different models mentioned above are tested using both the CFI dataset and the Embedding on the Wall dataset to assess them under different circumstances. As a pre-processing step, the faces from the CFI dataset must be extracted with the face detection model with the highest performance as explained in section 4.1. Furthermore, to ensure the algorithms provide better solutions to the clustering problem, it is important to reduce the dimensionality of the face embeddings before applying the clustering algorithms (Paukkeri et al., 2011; Kaski, 1998). Additionally, Ruyters has shown that t-distributed Stochastic Neighbour Embedding (t-SNE) can be used to create a clear visualization where the clusters are more separable rather than using Principal Components Analysis (PCA) to reduce the embeddings' dimension as can be seen in Figure 17. Therefore, after calculating the vector representation of all faces, they are reduced using t-SNE.

Both clustering algorithms are tested using different parameter input values, together with all seven t-SNE reduced face embeddings. The following evaluation metrics were applied to assess the model's performance:

- **Adjusted Rand Index (ARI):** The ARI-score is used for calculating the similarity between two clusterings. For this thesis, it can be used to describe how similar the clustering from the clustering algorithms is to the actual clustering as defined by each dataset. Furthermore, the ARI-score is adjusted to take into account that an algorithm can 'guess' a cluster correctly. The equation for the ARI-score is shown in equation 6.

Let $F = [f_1, f_2, \dots, f_n]$ be a sorted list of face embeddings.

Let $C = [c_1, c_2, \dots, c_n]$ be the sorted list of predicted cluster labels.

Let $G = [g_1, g_2, \dots, g_n]$ be the sorted list of ground-truth cluster labels.

The lists are sorted such that for every $i \in \{1, \dots, n\}$ the pair of f_i is denoted by (c_i, g_i) . A pair (c_i, g_i) is said to *agree* if and only if $c_i = g_i$.

Then the ARI-score is defined as:

$$\text{RI} = \frac{\text{\#agreeing_pairs}}{\text{\#pairs}} \quad (5)$$

$$\text{ARI} = \frac{(\text{RI} - \mathbb{E}[\text{RI}])}{(\text{max}(\text{RI}) - \mathbb{E}[\text{RI}])} \quad (6)$$

- **Relative number of clusters:** This shows how close the predicted number of clusters is to the actual number of clusters, useful to understand the minimum number of operation (merging or splitting clusters) needed to transform the predicted clustering into the actual clustering. The calculation is shown in Equation 7.

$$\text{rel} = \text{\#predicted_clusters} - \text{\#actual_clusters} \quad (7)$$

5 Interactive Visualization

The final component of this thesis is the development of an interactive visualization, allowing users to interact with the models by exploring, verifying, and adjusting the results of the face detection and recognition models. Moreover, this visualization was designed to broaden the scope from political image analysis to all datasets with a similar structure as the Tool-Test dataset. This section explains the functionalities and design choices embedded in the tool.

5.1 User Tasks

A potential user for this tool would be a researcher who is in the possession of a dataset similar to the Tool-Test dataset and is interested in common co-occurring persons in images or their corresponding attributes. To satisfy the interest of this user, a list of user tasks in accordance with the methodology described in (Munzner, 2014) has been defined as shown in Table 1.

Level	Action	Target	Task description (participant)	Realization
Analyze	Annotate	Features	Be able to add, remove, and edit any faces in an image uploaded by the user when creating a new database	Interactive visualization
			Be able to add, remove, and edit any cluster in the clustering when creating a new database	
Search	Browse	Features	Be exposed to all images uploaded by the user when creating a new database, in order to become aware of what the detection model lacks in	Allow user to scroll through all images with their detections
	Explore	Features	Be exposed to all face embeddings when creating a new database and the clustering algorithm’s parameters, in order to become aware of how these parameters affect the clustering	Allow user to play with the parameters and show them the resulting clustering
	Lookup	Dependency	Find out who co-occurs on an image in which the person of interest exists	Horizontal or vertical bar chart (ordered)

			Find out what values of some additional data are related to the images that the person of interest (and a second selected person) is in	
		Extremes	Find out which person co-occurs the most with the person of interest in a certain database	
			Find out which value, of all values of some additional data related to the images that the person of interest (and a second selected person) is in, is most common	
	Locate	Features	Find out whether a specific person of interest exists in the clustering of a certain database	Face recognition model on user input
Query	Identify	Features	Find out in what image the person of interest (together with a second selected person) exist in	Scroll through all images with selected person(s)

Table 1: The task abstraction for the interactive visualization, based on the methodology described in Munzner (2014)

5.2 Design Sketch

In accordance with the user tasks, a sketch of the tool is provided in Figure 7 of which all functions are explained in this section.




Figure 7: Sketch of the interactive visualization, consisting of two pages.

5.2.1 The 'Visualization' Page

In Figure 7a, a sketch of the *Visualization* page of the tool is shown. Here the user is able to select one of the databases on which a simple face recognition model is trained as explained by Ruyters, and the user can input an image from their pc. After which they can select a face detected on the image or draw a bounding box themselves. The face recognition model then tries to predict the correct cluster based on the selected face, an example of the predicted cluster is shown to check for the user manually. If the user already knows who he is interested in, he can select any cluster from a radio-menu instead of inputting an image.

The *explore* section is then used to highlight the cluster as explained in section 5.3.1. Additionally, this section allows the user to scroll through all images in which the selected person is detected. Since there are likely more images than can fit on the average screen, only a subset of the images is loaded and buttons allow the user to scroll through the images. The tool contains multiple of such image boxes, all having similar approaches as this one.

The section below the *explore* section shows all results, containing both the co-occurring persons and the additional information. In the first part, a visualization as described in

section 5.3.2 is used to display all persons who co-occur with the selected person of interest. Moreover, the user can select such co-occurring person to display all images in which both selected persons occur, including the bounding boxes for both these persons in distinct colours. To provide more context to these images, the user can click on each image to follow a URL, if provided by the dataset. At the bottom of this page, another visualization as explained in section 5.3.2 is used to display the additional data if provided by the dataset.

5.2.2 The 'New Data' Page

In Figure 7b, a sketch of the *New Data* page is shown. This is where the user can upload a new dataset which can either be a set of one or more images, or a CSV-file where the user can input additional information and URL's. An example of such a dataset is the Tool-Test dataset.

After all data is processed and the face detection model is finished detecting all faces, the first few images are shown next to each other with the resulting bounding boxes plotted on them as explained in section 5.3.3. Within this *Face Detection* section, the user can adjust which detected bounding boxes should be kept. Moreover, the user can add new detection by manually drawing a box in the image. Additionally, the user can limit the size of images to remove potential noise, i.e. detections with unrecognizable faces due to small sizes and bad quality of the detected face. To address the problem of having many faces in an image, the user can show the face numbers to know which face to edit by clicking on a button.

After the user is satisfied with the detected faces, the user can continue to the *Face Clustering* section where the user can adjust the clusters resulting from the face clustering model. On the left side, a plot displays the different clusters as described in section 5.3.1. Moreover, the user can create a new cluster by drawing one in this visualization. On the right side, the user can initially adjust the parameters for the clustering algorithm to optimize them to their liking. In addition, the user can set a minimum size for detected faces to be considered for clustering. Therefore, allowing the user to remove potential noise affecting the resulting clustering's quality significantly. Besides, the user can select any cluster of which some example images will be displayed.

Once satisfied, the user can continue to the next phase allowing the user to rename a cluster, merge two clusters, or move a single image from one cluster to another. Moreover, the user can use the figure on the left to add a new cluster by drawing a boundary. After the user has finished clustering, the user can save the results to the database and use the visualization page to get the results on the new database.

5.3 Design Choices - Visuals

Different visual design principles must be applied for different visuals in different circumstances. This section serves to explain what visuals are used in what context, and what design principles were applied.

5.3.1 Clustering Visuals

The clusterings resulting from the face clustering model can be displayed using a scatter plot, where each face is denoted by an x- and y-coordinate denoted by the corresponding t-SNE reduced embedding. For the *Face Clustering* section, there can be many more clusters than unique colours distinguishable by a human, which is 12 colours according to Munzner (2014). However, if there are more clusters than colours, it leads to colours being reused. Therefore a larger colour set is needed and thus a set of 20 unique colours¹¹ is used. If a user is still unsure what colour belongs to what cluster, he can enable and disable the cluster in the legend of the scatter plot to see where the dots dis-/reappear. An example of such a figure can be found in Figure 8.

A similar approach is used for the *explore* section, but only using two distinct colours. Therefore, the cluster in which the selected person of interest is recognized will be highlighted by colouring this cluster red within the scatter plot where all other points are light-gray.




Figure 8: Example image of a scatter plot used for the 'Face Clustering' section.

5.3.2 Co-occurring Persons & Additional Information Visuals

Visualizing the people who co-occur with the selected person of interest can be done by using a horizontal bar chart where the y-axis encodes the co-occurring person, and the x-axis is the length of a bar encoding the number of times that person co-occurs with the person of interest in an image. Similarly, the additional data, if provided, is shown to the user via a vertical bar chart. The x-axis describes all values for the selected additional data, and the y-axis is the length of a bar encoding the number of times a certain value occurs

¹¹<https://cmap-docs.readthedocs.io/en/latest/catalog/qualitative/seaborn:tab20/>

in all images where the person of interest occurs in, and if selected, the person belonging to a selected bar occurs in. Moreover, both figures are sorted in descending order based on the count, allowing the user to easily find the most common value. An example of the latter bar chart is shown in Figure 9




Figure 9: Example image of a bar chart for some additional data named 'category'.

5.3.3 Bounding Boxes

The bounding boxes must be easily distinguished from the image it is in. Since every image can have different colours, a colour is used that rarely occurs in the real world while still being distinguishable in most circumstances, lime-green. An example image is shown in Figure 10.

For the *Visualization* page, two colours are needed to encode both the selected person of interest and a selected co-occurring person. Since the two bounding boxes must be distinguished from each other and the image itself, the colours lime-green and magenta are used.



Figure 10: Example image of the plotted bounding boxes on an image.

6 Results

6.1 Face Detection Model Selection

In Table 2, the results from all different YOLOv12 versions are shown. Here, it is shown that YOLOv12s scores best on almost all evaluation metrics and has the lowest training time which is more ideal for training the larger YOLOv12 model, as explained in section 4.1.1, since it requires more data and more epochs. Therefore, the larger YOLO model, from now on referred to as YOLOv12s-face, is trained using the YOLOv12s architecture.

model version	YOLOv12n	YOLOv12s	YOLOv12m	YOLOv12l	YOLOv12x
mAP50	0.495	0.506	0.494	0.464	0.338
mAP50-95	0.265	0.279	0.267	0.249	0.180
precision	0.760	0.762	0.731	0.770	0.655
recall	0.433	0.437	0.433	0.390	0.290
runtime (mm:ss)	05:20	05:02	08:02	14:39	32:00

Table 2: Evaluation metrics for different versions of the YOLOv12 architecture trained using the FDD-WF small subset dataset.

Since the YOLOv12s-face model is trained for 100 epochs, there exists a chance that it may overfit to the training or validation data. To establish whether and when the model is overfitting, the losses are compared for the training and validation set as shown in Figure 11. From the figure, it is clear that the validation loss stops reducing around epoch 60 and even starts to slowly increase later on while the training loss is still steadily decreasing. This is a sign of overfitting and therefore the model weights for epoch 60 are used as the final YOLOv12s-face version.



Figure 11: Df1-loss comparison between the training and validation phase while training the YOLOv12s model.

The final results of the different face detection models are shown in Table 3. It is shown that for all models, there is a significant increase in the overall intersection over union and the accuracy compared with the base model. Among these models, YOLOv12s-face scores the best according to these evaluation metrics and YOLOv6n-face is the fastest. Even though the YOLOv6n-face model is approximately 17% faster than the YOLOv12s-face model, it is only 15ms faster per image ($\tilde{2.5}$ images per second faster). For the purpose of the interactive visualization, more time will be saved by using a more accurate and slower model, than by using a faster and less accurate model since less manual modification are needed. Therefore, YOLOv12s-face is for our purposes the best face detection model among the ones tested.

model	IoU (entire image)	Accuracy (IoU $\geq$ 0.5)	Accuracy (IoU $\geq$ 0.7)	Latency (s)
(base model)	0,473	0,561	0,224	0,455
InsightFace	0,821	0,877	0,775	0,086
YOLOv6n-face	0,830	0,875	0,770	0,071
YOLOv11l-face	0,825	0,870	0,766	0,222
YOLOv12s-face	0,843	0,883	0,789	0,086

Table 3: Evaluation metrics for face detection models tested using the FD-test dataset.

6.2 Face Clustering

In Figures 18-24, the t-SNE reduced embeddings are plotted for each model and colored as to the actual clusters. For both the CFI dataset and the Embedding on the Wall dataset, the faces embedded using DeepFace and OpenFace are all over the place as if the data would be considered noise. Since there are no visual clusters for these two models, it is not possible for a distance-based clustering algorithm to find clusters close to the actual clusters of these datasets. Therefore these two models will not be further considered for embedding faces. For all other models, there are visual clusters, so these models will be considered for the next step. But it must be noted that there always seems to be at least one cluster of noise which makes it harder for distance-based clustering algorithms to create clusters close to the actual clusters.

In Tables 5 & 6, the resulting evaluation metrics are shown for different values of the epsilon parameter used for the DBSCAN algorithm for both datasets. Both tables show that the ArcFace model and the Dlib model have a decent ARI-score, but have significantly more predicted clusters than there actually are. For the CFI dataset, the FaceNet and face_recognition models both have high ARI-scores while remaining close to the actual number of clusters. However, looking at the results using the Embedding on the Wall dataset, the embedding by FaceNet shows a larger difference in the number of clusters than there actually are, the same goes for the VGG_Face model. For both datasets, it holds that using the face_recognition model to embed faces leads to the clustering which is most similar to the actual clusters when using the DBSCAN algorithm. Moreover, it is notable

that the range of optimal values for the epsilon parameter differs for both datasets.

In Tables 7 & 8, the resulting evaluation metrics are shown for different values of the threshold for the Chinese Whispers clustering algorithm for both datasets. For the CFI dataset, all models used for calculating the embeddings perform well with both a high ARI-score and with the predicted number of clusters being close to the actual number of clusters. The face_recognition model performs best for this dataset with the highest ARI-score. For the Embedding on the Wall dataset, the face_recognition model and VGG_Face model perform best, of which the face_recognition model results in the highest ARI-score. But for both models, the resulting number of clusters is further from the actual number of clusters than it was for the CFI dataset.

To compare the two algorithms with each other, the best results of both models are shown in Table 4. Overall, the ARI-scores for both clustering algorithms are good, but DBSCAN has the overall highest scores. Moreover, the time to compute the clusters grows significantly more for the Chinese Whispers algorithm than the DBSCAN algorithm when the input size is increased.

model	CFI dataset			Embedding on the Wall		
	ARI	rel	latency (s)	ARI	rel	latency (s)
DBSCAN	0.950	1	0.060	0.964	0	0.022
Chinese Whispers	0.950	1	9.515	0.905	-3	0.626

Table 4: Best results for both clustering algorithms in terms of evaluation metrics where ARI = adjusted rand score and rel = the difference in resulting and actual clusters.

6.3 Interactive Visualization

To develop a better understanding of the functionalities of the interactive visualization, a few use cases are provided.

6.3.1 Use Case 1

Suppose you are a researcher interested in the politicians represented in the Dutch news header images during the last elections. And you are especially interested in the people who are commonly shown in these header images together with an unknown person of whom you have an image. Moreover, you are interested in what kind of issues those articles discuss. Now, this interest overlaps with the interest of Ruyters, of which the database is already available when you navigate to the *Visualization* page and select the database *Embedding-On-The-Wall-fixed.db*, which yields from using the *New Data* page on the Tool-Test dataset. Now, the user can select any person of interest from the radio menu, but this user does not know the name of the person. Therefore, the user can upload an image and select the person he is interested in. As can be seen in Figure 12, the YOLOv12s model detects three faces in the uploaded image. To find out which face matches which name in the radio menu, the user can check the *show face numbers* checkbox. When selecting a

face in the radio menu, a simple face recognition model predicts that the face matches the face of Geert Wilders. The user can now check if the two faces shown on the right-hand side match, and if they do not match, the user can select one of the clusters from the radio menu and try to find a matching cluster to the uploaded image. Note that it can be the case that such a cluster does not exist in the database.




Figure 12: Section on the *visualization* page where the user can select a person of interest, displaying that there are three faces detected (two in front, one in the back) and that the face recognition model correctly identifies Geert Wilders as the selected person of interest.

Continuing to the results via clicking the *Show Results* button yields two new sections: the *Explore Cluster* section and the *Results* section. In the *Explore Cluster* section, the user can check whether all images indeed include Geert Wilders. In the *Results* section, the user can find what he is interested in. Here, the user finds all persons who are found in the same images as Geert Wilders, and by selecting a person in the bar chart, the user can check all images to see if both persons indeed exist in all of them. Moreover, the user can select any additional information of his interest, which is *Issues_paragraphs* in this case. Then, the user is able to see the kind of issues discussed in news articles related to the two selected persons as can be seen in Figure 13. Additionally, the user can click on any image on this page to navigate to the corresponding news article if needed.




Figure 13: The *Explore Cluster & Results* sections on the *visualization* page where the user can explore the cluster containing the person of interest, select another person who co-occurs in at least one image with the person of interest, check those corresponding images and see which issues are commonly discussed in those news articles.

The results show that Geert Wilders occurs 65 times in the loaded dataset, of which he occurs 14 times with Pieter Omtzigt. Moreover, migration is the most common issue discussed in news articles with Geert Wilders and Pieter Omtzigt in the picture lead.

6.3.2 Use Case 2

Consider a user who is interested in celebrities and whether there are any celebrities who are commonly in the same images in a dataset the user gathered. Since there is no database by default in the tool, the user must create one. this can be done by navigating to the *New Data* page. Here the user uploads an arbitrary amount of images in jpg or png format from his dataset. Once the face detection model is done with detecting all faces, a new section appears where the user can add and remove any face from the set of selected faces as shown in Figure 14.




Figure 14: *Face Detection* section on the *New Data* page where a user can add and remove any face from the set of selected faces. Additionally, any potential noise (background images) can be removed by setting a minimum size for faces.

Once the user is satisfied with the set of selected faces, he proceeds to the *Face Clustering* section where the parameters for the DBSCAN clustering algorithm can be optimized. The progression of this is shown in Figure 15. Note that some faces have gone to the unknown class since they do not meet the requirement of having 8 neighbours within 2.6 distance.




Figure 15: *Face Clustering* section on the *New Data* page in the visualization, showing a dataset before (upper) and after (lower) fine-tuning the parameters for the DBSCAN clustering algorithm.

When the user finds values for the parameters satisfying their expectations, the user can start editing clusters by adding and merging clusters, and by moving single images from one cluster to another. Moreover, the user can name the clusters to allow for easier interpretation when going through the results. This process is shown in Figure 16.




Figure 16: *Face Clustering* section on the *New Data* page in the visualization, where users can add, edit, and remove any cluster to their liking.

When the user is done with the manual clustering, he can save the database with an inputted name. To see the results and check who are commonly in the same image, the user can follow similar steps as in Use Case 1 by selecting the recently saved database.

The difference is that the user now knows which person(s) he is interested in so, instead of inputting an image, he can just select a cluster from the radio menu. Besides, the user did only input images and not a CSV-file with additional data. Therefore, no additional information as shown in the lower image in Figure 13 will be displayed.

7 Conclusion and Discussion

This thesis is set out to improve the face detection and clustering models based on the findings of Ruyters. Moreover, we developed an interactive visualization to use those models and create an understanding of their functioning for a broader scope than Ruyter’s interest. This section explains the key findings, limitations, and future work for each of the models and the developed tool.

7.1 Face Detection

For face detection, it is recommended to use the YOLOv12s-face model since it performs best among the models tested according to metrics used during evaluation. However, there is still a problem with this model: sometimes the model predicts multiple faces on one face which have therefore overlapping bounding boxes. On the other side, when two faces are behind each other, the model can predict both faces correctly but they still have overlapping boxes. Duplicated boxes can, in the worst case, lead to wrong conclusions when using the visualization. For example, if some person has many duplicates, they may be overrepresented compared to what they actually are.

Future improvements could be trying to solve the problem of duplicates, and using larger datasets to train the model on, which may lead to higher IoU-scores and higher accuracies.

7.2 Face Clustering

Regarding face clustering, there are two algorithms tested in this thesis. Both algorithms perform best when used with faces embedded by the face_recognition model. Moreover, both models have high ARI-scores when used with the correct parameters. However, finding the correct values for these parameters can be time-consuming. If you prefer an easy setup and do not care much about runtime, the Chinese Whispers algorithm is a great option. On the other hand, if the speed of the algorithm is of greater importance, the DBSCAN clustering algorithm is recommended. Furthermore, noise could be problematic for the performance of a clustering algorithm.

Analysis of other methods for vectorization of faces may minimize the latter issue. However, using a different clustering algorithm could improve both problems, think of algorithms like Clustering Using REpresentatives (CURE) (Guha et al., 2001).

7.3 Interactive Visualization

Some problems with the visualization may occur when trying to cluster datasets that contain small, low-quality images or many single-occurring faces. Both tested clustering algorithms are based on distances between two points and therefore lack the ability to handle this kind of noise correctly, which then must be corrected by the user manually. This can partially be solved within the visualization. For example, if the user would be

allowed to select some clusters that form unique faces, based on which a face recognition model would be trained and after which it tries to predict all remaining faces to one of those clusters or an unknown cluster. It could potentially take away a lot of manual labelling by the user, if implemented correctly and user-friendly.

More improvements could be implemented for the face recognition model used when identifying which cluster belongs to the selected face. This thesis implemented the same pipeline as used by Ruyters. However, this is far from good and is inconsistent. Analyzing different face recognition models would be necessary to partially solve this problem.

Further expansion of the tool may include expansion of the additional information section and identifying other objects or patterns in the uploaded images to provide more insights of different aspects to the user:

- **Additional information:** This section could be expanded in different ways. One could be to allow different data types to be entered instead of lists of items only. For example, allowing numbers could introduce extra insights via a histogram. Moreover, time could introduce a timestamp option, allowing the user to focus on a specific period. In addition, other graphs, like a multi-bar chart, could introduce the option to directly compare two persons with each other.
- **Object detection & pattern recognition:** Via object detection and pattern recognition, new insights can be introduced by finding common items. This could extend the scope of this tool even further than just looking at co-occurring people in images.

References

- Amel Ibrahim Al Ali, Sheeja Rani S, and Ahmed M. Khedr. Enhancing financial distress prediction through integrated chinese whisper clustering and federated learning. *Journal of Open Innovation: Technology, Market, and Complexity*, 10(3):100344, 2024. ISSN 2199-8531. doi: 10.1016/j.joitmc.2024.100344.
- Momina Liaqat Ali and Zhou Zhang. The yolo framework: A comprehensive review of evolution, applications, and benchmarks in object detection. *Computers*, 13(12), 2024. ISSN 2073-431X. doi: 10.3390/computers13120336. URL <https://www.mdpi.com/2073-431X/13/12/336>.
- BEP. Face detection dataset. <https://universe.roboflow.com/bep-jzm3u/face-detection-nmcfg>, mar 2025. URL <https://universe.roboflow.com/bep-jzm3u/face-detection-nmcfg>. visited on 2025-04-03.
- Chris Biemann. Chinese whispers - an efficient graph clustering algorithm and its application to natural language processing problems. In Rada Mihalcea and Dragomir Radev, editors, *Proceedings of TextGraphs: the First Workshop on Graph Based Methods for Natural Language Processing*, pages 73–80, New York City, June 2006. Association for Computational Linguistics. URL <https://aclanthology.org/W06-3812/>.
- Adam Geitgey. Face recognition: Recognize and manipulate faces from python or from the command line, 2020. URL <https://github.com/ageitgey/facerecognition>.
- Scott Grissom, Myles F. McNally, and Tom Naps. Algorithm visualization in cs education: comparing levels of student engagement. In *Proceedings of the 2003 ACM Symposium on Software Visualization*, SoftVis '03, page 87–94, New York, NY, USA, 2003. Association for Computing Machinery. ISBN 1581136420. doi: 10.1145/774833.774846.
- Sudipto Guha, Rajeev Rastogi, and Kyuseok Shim. Cure: an efficient clustering algorithm for large databases. *Information Systems*, 26(1):35–58, 2001. ISSN 0306-4379. doi: [https://doi.org/10.1016/S0306-4379\(01\)00008-4](https://doi.org/10.1016/S0306-4379(01)00008-4). URL <https://www.sciencedirect.com/science/article/pii/S0306437901000084>.
- Jia Guo, Jiankang Deng, Alexandros Lattas, and Stefanos Zafeiriou. Sample and computation redistribution for efficient face detection, 2021. URL <https://arxiv.org/abs/2105.04714>.
- Glenn Jocher, Ayush Chaurasia, and Jing Qiu. YOLO by Ultralytics, January 2023. URL <https://github.com/ultralytics/ultralytics>.
- S. Kaski. Dimensionality reduction by random mapping: fast similarity computation for clustering. In *1998 IEEE International Joint Conference on Neural Networks Proceedings. IEEE World Congress on Computational Intelligence (Cat. No.98CH36227)*, volume 1, pages 413–418 vol.1, 1998. doi: 10.1109/IJCNN.1998.682302.

- Zengmin Lin, Chaoqun Hong, Weiwei Zhuang, and Keshou Wu. Face clustering based on fusion of face tracking and optimization. In *Proceedings of the 2020 9th International Conference on Computing and Pattern Recognition*, ICCPR '20, page 208–212, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450387835. doi: 10.1145/3436369.3437420.
- Tamara Munzner. Visualization analysis and design, 2014.
- T.L. Naps. Jhave: supporting algorithm visualization. *IEEE Computer Graphics and Applications*, 25(5):49–55, 2005. doi: 10.1109/MCG.2005.110.
- Mari-Sanna Paukkeri, Ilkka Kivimäki, Santosh Tirunagari, Erkki Oja, and Timo Honkela. Effect of dimensionality reduction on different distance measures in document clustering. In Bao-Liang Lu, Liqing Zhang, and James Kwok, editors, *Neural Information Processing*, pages 167–176, Berlin, Heidelberg, 2011. Springer Berlin Heidelberg. ISBN 978-3-642-24965-5.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- Alina Petukhova, João P. Matos-Carvalho, and Nuno Fachada. Text clustering with large language model embeddings. *International Journal of Cognitive Computing in Engineering*, 6:100–108, December 2025. ISSN 2666-3074. doi: 10.1016/j.ijcce.2024.11.004. URL <http://dx.doi.org/10.1016/j.ijcce.2024.11.004>.
- Bernhard Preim, Renata Raidou, Noeska Smit, and Kai Lawonn. Chapter 9 - an introduction to visual analytics. In Bernhard Preim, Renata Raidou, Noeska Smit, and Kai Lawonn, editors, *Visualization, Visual Analytics and Virtual Reality in Medicine*, The MICCAI Society Book Series, pages 225–272. Academic Press, 2023. ISBN 978-0-12-822962-0. doi: <https://doi.org/10.1016/B978-0-12-822962-0.00018-3>. URL <https://www.sciencedirect.com/science/article/pii/B9780128229620000183>.
- Wies Ruyters. Embedding, embedding on the wall; who is the most prominent politician of them all? exploring automated methods to study multi-modal political news coverage. 2025.
- Purvi Saraiya, Clifford A. Shaffer, D. Scott McCrickard, and Chris North. Effective features of algorithm visualizations. *SIGCSE Bull.*, 36(1):382–386, March 2004. ISSN 0097-8418. doi: 10.1145/1028174.971432.
- Florian Schroff, Dmitry Kalenichenko, and James Philbin. Facenet: A unified embedding for face recognition and clustering. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June

2015. URL https://www.cv-foundation.org/openaccess/content_cvpr_2015/html/Schroff_FaceNet_A_Unified_2015_CVPR_paper.html.
- Sefik Ilkin Serengil and Alper Ozpinar. Lightface: A hybrid deep face recognition framework. In *2020 Innovations in Intelligent Systems and Applications Conference (ASYU)*, pages 23–27. IEEE, 2020. doi: 10.1109/ASYU50717.2020.9259802. URL <https://ieeexplore.ieee.org/document/9259802>.
- Sefik Ilkin Serengil and Alper Ozpinar. Hyperextended lightface: A facial attribute analysis framework. In *2021 International Conference on Engineering and Emerging Technologies (ICEET)*, pages 1–4, 2021. doi: 10.1109/ICEET53442.2021.9659697.
- Sefik Ilkin Serengil and Alper Ozpinar. A benchmark of facial recognition pipelines and co-usability performances of modules. *Bilisim Teknolojileri Dergisi*, 17(2):95–107, 2024. doi: 10.17671/gazibtd.1399077. URL <https://dergipark.org.tr/en/pub/gazibtd/issue/84331/1399077>.
- Karen Shiells and Peter Pham. Unsupervised clustering for language identification, 2010. URL https://snap.stanford.edu/class/cs224w-2010/proj2010/28_writeup.pdf.
- Siffi Singh and Hsueh-Ming Hang. A continuous learning system for face clustering and recognition. In *2021 IEEE International Conference on Consumer Electronics (ICCE)*, pages 1–6, 2021. doi: 10.1109/ICCE50685.2021.9427771.
- Jacob Solawetz. What is mean average precision (map) in object detection?, May 2020. URL <https://blog.roboflow.com/mean-average-precision/>.
- James J Thomas and Kristin A Cook. *Illuminating the Path: The Research and Development Agenda for Visual Analytics*. IEEE Computer Society, Los Alamitos, CA, United States(US)., 05 2005. URL <https://www.osti.gov/biblio/912515>.

A Appendix

A.1 Pseudocode

Algorithm 1 IoU-based Accuracy Calculation

Require: B : list of predicted bounding boxes

Require: A : list of actual bounding boxes

Require: $iou_threshold$: threshold for IoU

```
1: Initialize  $tp \leftarrow 0$ ,  $fp \leftarrow 0$ ,  $fn \leftarrow 0$ 
2: for each  $box$  in  $B$  do
3:    $iou\_scores \leftarrow []$ 
4:    $max\_iou \leftarrow -1$ 
5:   for each  $bb$  in  $A$  do
6:     Append IoU between  $box$  and  $bb$  to  $iou\_scores$ 
7:   end for
8:   if  $iou\_scores \neq []$  then
9:      $max\_iou \leftarrow \max(iou\_scores)$ 
10:  end if
11:  if  $max\_iou \geq iou\_threshold$  then
12:     $tp \leftarrow tp + 1$ 
13:    Remove box at index  $\text{argmax}(iou\_scores)$  from  $A$ 
14:  else
15:     $fp \leftarrow fp + 1$ 
16:  end if
17: end for
18:  $fn \leftarrow fn + \text{length of } A$ 
19: Accuracy =  $\frac{tp}{tp+fp+fn}$ 
```

A.2 Dimensionality Reduced Clustering




Figure 17: Visualization of embeddings clusters where they are reduced in dimensionality using t-SNE (a) and PCA (b) as shown in the paper from Ruyters.

A.3 t-SNE Reduced Embeddings with Different Models




Figure 18: Actual cluster of all faces in both the CFI dataset (left) and the Embedding on the Wall dataset (right) embedded using the ArcFace model. Each color defines a unique person.




Figure 19: Actual cluster of all faces in both the CFI dataset (left) and the Embedding on the Wall dataset (right) embedded using the DeepFace model. Each color defines a unique person.




Figure 20: Actual cluster of all faces in both the CFI dataset (left) and the Embedding on the Wall dataset (right) embedded using the Dlib model. Each color defines a unique person.




Figure 21: Actual cluster of all faces in both the CFI dataset (left) and the Embedding on the Wall dataset (right) embedded using the face_recognition model. Each color defines a unique person.




Figure 22: Actual cluster of all faces in both the CFI dataset (left) and the Embedding on the Wall dataset (right) embedded using the FaceNet model. Each color defines a unique person.




Figure 23: Actual cluster of all faces in both the CFI dataset (left) and the Embedding on the Wall dataset (right) embedded using the OpenFace model. Each color defines a unique person.




Figure 24: Actual cluster of all faces in both the CFI dataset (left) and the Embedding on the Wall dataset (right) embedded using the VGG_Face model. Each color defines a unique person.

A.4 Testing Parameters for Clustering with DBSCAN

params	ArcFace		Dlib		FaceNet		face_recognition		VGG_Face	
eps	ARI	rel	ARI	rel	ARI	rel	ARI	rel	ARI	rel
0.1	0.003	1616	0.003	1630	0.004	1588	0.003	1605	0.003	1599
0.2	0.006	1477	0.005	1521	0.007	1449	0.007	1441	0.006	1476
0.3	0.010	1347	0.008	1398	0.011	1289	0.012	1291	0.009	1359
0.4	0.015	1175	0.012	1249	0.020	1086	0.017	1136	0.014	1201
0.5	0.024	997	0.019	1098	0.039	876	0.030	943	0.021	1049
0.6	0.047	799	0.037	924	0.084	683	0.063	735	0.030	906
0.7	0.104	619	0.082	723	0.153	516	0.156	537	0.053	742
0.8	0.226	431	0.170	547	0.305	376	0.358	368	0.119	570
0.9	0.384	311	0.298	423	0.426	279	0.577	245	0.221	434
1.0	0.557	218	0.430	335	0.547	207	0.659	173	0.356	323
1.1	0.677	150	0.507	262	0.636	154	0.726	132	0.466	252
1.2	0.729	107	0.640	194	0.708	115	0.789	91	0.543	197
1.3	0.754	85	0.697	155	0.735	94	0.821	67	0.636	155
1.4	0.795	63	0.751	124	0.774	72	0.854	52	0.697	120
1.5	0.807	51	0.768	103	0.811	60	0.870	43	0.756	90
1.6	0.807	41	0.785	87	0.823	51	0.882	37	0.817	69
1.7	0.831	30	0.801	76	0.843	45	0.891	31	0.850	52
1.8	0.870	21	0.807	65	0.849	34	0.901	22	0.874	39
1.9	0.874	17	0.810	57	0.872	26	0.916	15	0.887	31
2.0	0.874	15	0.812	53	0.882	22	0.931	10	0.900	26
2.1	0.811	13	0.820	45	0.890	16	0.938	8	0.907	20
2.2	0.816	10	0.828	39	0.895	15	0.938	8	0.914	15
2.3	0.808	8	0.830	34	0.896	13	0.941	7	0.914	15
2.4	0.806	6	0.806	26	0.896	12	0.941	7	0.921	13
2.5	0.808	5	0.814	19	0.901	10	0.941	7	0.924	11
2.6	0.808	5	0.774	14	0.904	8	0.946	6	0.924	11
2.7	0.808	4	0.774	14	0.904	8	0.947	5	0.928	9
2.8	0.808	4	0.774	14	0.909	6	0.947	5	0.928	9
2.9	0.810	3	0.717	10	0.922	5	0.949	3	0.930	7
3.0	0.810	3	0.722	8	0.922	5	0.949	3	0.934	5
3.1	0.809	2	0.722	7	0.921	4	0.949	3	0.934	5
3.2	0.809	2	0.724	4	0.921	4	0.950	2	0.939	4
3.3	0.809	2	0.724	4	0.921	4	0.950	2	0.939	4
3.4	0.809	2	0.699	3	0.921	4	0.950	2	0.939	4
3.5	0.812	1	0.699	3	0.921	4	0.950	2	0.939	4
3.6	0.812	1	0.700	1	0.921	4	0.950	2	0.939	4
3.7	0.812	1	0.700	1	0.921	4	0.950	2	0.939	4
3.8	0.719	0	0.700	1	0.921	4	0.950	2	0.939	4

3.9	0.719	0	0.673	0	0.921	4	0.950	2	0.939	4
4.0	0.719	0	0.673	0	0.921	4	0.950	2	0.939	4
4.1	0.623	-1	0.673	0	0.921	4	0.950	2	0.939	4
4.2	0.623	-1	0.672	-1	0.921	4	0.950	2	0.894	3
4.3	0.623	-1	0.673	-2	0.922	3	0.950	2	0.894	3
4.4	0.623	-1	0.673	-2	0.924	2	0.950	2	0.894	3
4.5	0.623	-1	0.673	-2	0.924	2	0.950	2	0.894	2
4.6	0.623	-1	0.673	-2	0.924	2	0.950	2	0.894	2
4.7	0.623	-1	0.672	-3	0.924	2	0.950	2	0.894	2
4.8	0.604	-2	0.577	-4	0.924	2	0.950	2	0.894	2
4.9	0.603	-3	0.577	-4	0.924	2	0.950	2	0.747	0
5.0	0.603	-3	0.401	-5	0.886	1	0.950	2	0.747	0
5.1	0.603	-3	0.401	-5	0.886	1	0.950	2	0.747	0
5.2	0.603	-3	0.292	-6	0.886	1	0.950	2	0.747	0
5.3	0.603	-3	0.292	-6	0.883	0	0.950	1	0.747	0
5.4	0.603	-3	0.292	-6	0.883	0	0.950	1	0.747	0
5.5	0.603	-3	0.292	-6	0.883	0	0.950	1	0.747	0
5.6	0.603	-3	0.292	-6	0.883	0	0.950	1	0.747	0
5.7	0.603	-3	0.292	-6	0.883	0	0.950	1	0.714	-1
5.8	0.603	-3	0.292	-6	0.883	0	0.950	1	0.714	-1
5.9	0.603	-3	0.292	-6	0.883	0	0.950	1	0.715	-2
6.0	0.510	-4	0.292	-7	0.883	0	0.904	0	0.715	-2
6.1	0.510	-4	0.292	-7	0.883	0	0.904	0	0.715	-2
6.2	0.510	-4	0.292	-7	0.883	0	0.904	0	0.715	-2
6.3	0.510	-4	0.286	-8	0.883	0	0.904	0	0.715	-2
6.4	0.510	-4	0.286	-8	0.883	0	0.904	0	0.715	-2
6.5	0.510	-4	0.286	-8	0.883	0	0.904	0	0.715	-2
6.6	0.510	-4	0.286	-8	0.883	0	0.904	0	0.715	-2
6.7	0.510	-4	0.279	-9	0.883	0	0.904	0	0.685	-3
6.8	0.510	-4	0.279	-9	0.883	0	0.904	0	0.685	-3
6.9	0.510	-4	0.189	-10	0.847	-1	0.904	0	0.655	-4
7.0	0.510	-4	0.189	-10	0.847	-1	0.904	0	0.655	-4
7.1	0.510	-4	0.189	-10	0.847	-1	0.904	0	0.655	-4
7.2	0.510	-4	0.118	-11	0.847	-1	0.904	0	0.655	-4
7.3	0.510	-4	0.118	-11	0.847	-1	0.904	0	0.655	-4
7.4	0.510	-4	0.118	-11	0.847	-1	0.904	0	0.655	-4
7.5	0.510	-4	0.118	-11	0.847	-1	0.882	-1	0.655	-4
7.6	0.510	-4	0.067	-12	0.847	-1	0.882	-1	0.655	-4
7.7	0.510	-4	0.067	-12	0.847	-1	0.882	-1	0.596	-5
7.8	0.510	-4	0.067	-12	0.847	-1	0.882	-1	0.519	-6
7.9	0.510	-4	0.067	-12	0.847	-1	0.882	-1	0.459	-7
8.0	0.510	-4	0.067	-12	0.847	-1	0.882	-1	0.459	-7
8.1	0.510	-4	0.066	-13	0.847	-1	0.882	-1	0.459	-7

8.2	0.510	-4	0.028	-14	0.847	-1	0.824	-2	0.459	-7
8.3	0.510	-4	0.028	-14	0.847	-1	0.824	-2	0.459	-7
8.4	0.510	-4	0.028	-14	0.780	-2	0.824	-2	0.459	-7
8.5	0.510	-4	0.028	-14	0.623	-3	0.824	-2	0.459	-7
8.6	0.510	-4	0.028	-14	0.623	-3	0.824	-2	0.459	-7
8.7	0.370	-5	0.028	-14	0.602	-4	0.824	-2	0.459	-7
8.8	0.370	-5	0.028	-14	0.561	-5	0.824	-2	0.459	-7
8.9	0.370	-5	0.028	-14	0.561	-5	0.824	-2	0.459	-7
9.0	0.370	-5	0.028	-14	0.561	-5	0.824	-2	0.459	-7
9.1	0.370	-5	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.2	0.370	-5	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.3	0.370	-5	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.4	0.370	-5	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.5	0.370	-5	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.6	0.252	-6	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.7	0.252	-6	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.8	0.201	-7	0.028	-14	0.473	-6	0.741	-3	0.459	-7
9.9	0.201	-7	0.028	-14	0.200	-8	0.741	-3	0.459	-7
10.0	0.201	-7	0.028	-14	0.200	-8	0.741	-3	0.459	-7

Table 5: All tested values for the epsilon parameter for the DBSCAN algorithm using the CFI dataset where ARI = 'adjusted rand score' and rel = the relative number of clusters.

params	ArcFace		Dlib		FaceNet		face_recognition		VGG_Face	
eps	ARI	rel	ARI	rel	ARI	rel	ARI	rel	ARI	rel
0.1	0.005	326	0.004	329	0.006	320	0.007	315	0.006	318
0.2	0.008	312	0.006	318	0.010	305	0.013	293	0.009	307
0.3	0.010	302	0.009	307	0.012	295	0.019	275	0.011	296
0.4	0.011	297	0.010	301	0.015	284	0.022	261	0.013	290
0.5	0.014	290	0.012	294	0.023	262	0.030	238	0.015	281
0.6	0.015	285	0.014	288	0.030	242	0.043	211	0.019	267
0.7	0.019	272	0.016	277	0.039	217	0.083	172	0.022	256
0.8	0.023	261	0.020	264	0.063	185	0.118	147	0.024	248
0.9	0.027	246	0.027	244	0.098	153	0.242	108	0.036	222
1.0	0.036	226	0.039	222	0.207	123	0.371	82	0.054	186
1.1	0.049	199	0.053	205	0.283	91	0.495	49	0.079	162
1.2	0.070	174	0.096	166	0.373	76	0.671	37	0.110	142
1.3	0.101	142	0.163	141	0.488	59	0.750	27	0.169	121
1.4	0.139	114	0.215	117	0.539	44	0.782	20	0.235	100
1.5	0.198	93	0.272	90	0.675	29	0.820	15	0.327	76
1.6	0.325	73	0.427	67	0.699	23	0.904	11	0.383	64
1.7	0.386	61	0.522	51	0.725	19	0.925	8	0.507	51
1.8	0.479	46	0.593	42	0.761	12	0.934	7	0.570	44
1.9	0.527	37	0.671	31	0.800	6	0.934	7	0.653	36
2.0	0.572	30	0.694	27	0.793	3	0.934	7	0.686	30
2.1	0.662	20	0.743	23	0.672	-1	0.962	3	0.725	22
2.2	0.695	13	0.750	16	0.647	-4	0.964	1	0.791	15
2.3	0.736	9	0.750	16	0.647	-4	0.964	0	0.844	10
2.4	0.756	8	0.759	14	0.522	-5	0.964	0	0.743	7
2.5	0.709	4	0.734	7	0.503	-6	0.964	0	0.743	7
2.6	0.514	0	0.726	5	0.506	-7	0.956	-2	0.698	4
2.7	0.507	-2	0.658	1	0.506	-7	0.929	-3	0.668	-1
2.8	0.507	-2	0.633	0	0.506	-7	0.929	-4	0.668	-1
2.9	0.451	-4	0.514	-4	0.506	-7	0.929	-4	0.671	-2
3.0	0.451	-4	0.514	-4	0.499	-8	0.929	-4	0.671	-2
3.1	0.433	-5	0.508	-7	0.380	-9	0.929	-4	0.681	-3
3.2	0.353	-6	0.313	-9	0.380	-9	0.911	-5	0.408	-5
3.3	0.358	-7	0.249	-11	0.380	-9	0.849	-7	0.408	-5
3.4	0.179	-8	0.248	-12	0.361	-10	0.741	-9	0.400	-6
3.5	0.171	-9	0.232	-13	0.361	-10	0.741	-9	0.381	-7
3.6	0.171	-9	0.145	-14	0.368	-11	0.741	-9	0.366	-8
3.7	0.171	-9	0.145	-14	0.368	-11	0.741	-9	0.366	-8
3.8	0.171	-9	0.145	-14	0.352	-12	0.741	-9	0.364	-9
3.9	0.172	-10	0.004	-15	0.352	-12	0.690	-10	0.315	-10

4.0	0.138	-11	0.004	-15	0.352	-12	0.690	-10	0.315	-10
4.1	0.095	-12	0.004	-15	0.352	-12	0.690	-10	0.315	-10
4.2	0.097	-13	0.004	-15	0.352	-12	0.690	-10	0.313	-11
4.3	0.072	-14	0.004	-15	0.211	-13	0.654	-11	0.256	-12
4.4	0.072	-14	0.004	-15	0.211	-13	0.411	-12	0.256	-12
4.5	0.072	-14	0.004	-15	0.211	-13	0.411	-12	0.256	-12
4.6	0.072	-14	0.004	-15	0.211	-13	0.257	-13	0.212	-13
4.7	0.072	-14	0.004	-15	0.211	-13	0.257	-13	0.200	-14
4.8	0.072	-14	0.000	-16	0.211	-13	0.257	-13	0.200	-14
4.9	0.072	-14	0.000	-16	0.211	-13	0.257	-13	0.200	-14
5.0	0.072	-14	0.000	-16	0.211	-13	0.257	-13	0.200	-14
5.1	0.072	-14	0.000	-16	0.211	-13	0.257	-13	0.200	-14
5.2	0.072	-14	0.000	-16	0.211	-13	0.225	-14	0.200	-14
5.3	0.024	-15	0.000	-16	0.211	-13	0.192	-15	0.200	-14
5.4	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.200	-14
5.5	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.200	-14
5.6	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.200	-14
5.7	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.200	-14
5.8	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.200	-14
5.9	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.001	-15
6.0	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.001	-15
6.1	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.001	-15
6.2	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.001	-15
6.3	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
6.4	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
6.5	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
6.6	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
6.7	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
6.8	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
6.9	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.0	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.1	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.2	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.3	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.4	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.5	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.6	0.024	-15	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.7	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.8	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
7.9	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.0	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.1	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.2	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16

8.3	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.4	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.5	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.6	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.7	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.8	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
8.9	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.0	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.1	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.2	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.3	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.4	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.5	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.6	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.7	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.8	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
9.9	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16
10.0	0.000	-16	0.000	-16	0.153	-14	0.192	-15	0.000	-16

Table 6: All tested values for the epsilon parameter for DBSCAN using the Embedding on the Wall dataset where ARI = 'adjusted rand score' and rel = the relative number of clusters.

A.5 Testing Parameters for Clustering with Chinese Whispers

params threshold	ArcFace		Dlib		FaceNet		face_recognition		VGG_Face	
	ARI	rel	ARI	rel	ARI	rel	ARI	rel	ARI	rel
0.70	0.000	-16	0.092	-14	0.260	-13	0.000	-16	0.169	-14
0.71	0.000	-16	0.000	-16	0.105	-14	0.000	-16	0.170	-14
0.72	0.041	-15	0.102	-15	0.217	-13	0.000	-16	0.216	-14
0.73	0.000	-16	0.096	-14	0.000	-16	0.093	-15	0.218	-14
0.74	0.000	-16	0.102	-15	0.105	-14	0.116	-15	0.104	-14
0.75	0.041	-15	0.000	-16	0.136	-13	0.176	-14	0.209	-12
0.76	0.124	-12	0.204	-12	0.136	-13	0.029	-15	0.135	-14
0.77	0.220	-13	0.145	-13	0.136	-13	0.199	-13	0.271	-13
0.78	0.134	-12	0.316	-11	0.143	-12	0.204	-12	0.431	-10
0.79	0.134	-12	0.241	-11	0.143	-12	0.387	-10	0.252	-11
0.80	0.185	-10	0.193	-11	0.332	-10	0.256	-11	0.413	-9
0.81	0.403	-8	0.408	-9	0.462	-8	0.385	-9	0.428	-8
0.82	0.224	-8	0.245	-11	0.445	-7	0.603	-6	0.511	-7
0.83	0.277	-7	0.400	-9	0.749	-3	0.698	-4	0.628	-5
0.84	0.297	-6	0.526	-7	0.708	-3	0.674	-3	0.705	-4
0.85	0.368	-5	0.650	-3	0.812	-2	0.774	-2	0.697	-3
0.96	0.599	-3	0.650	-3	0.890	0	0.862	-1	0.850	-1
0.87	0.694	-2	0.671	-3	0.841	-1	0.863	-1	0.847	-1
0.88	0.793	0	0.730	-2	0.886	0	0.925	0	0.895	0
0.89	0.805	0	0.862	1	0.927	1	0.950	1	0.894	0
0.90	0.902	2	0.831	0	0.923	1	0.950	1	0.939	1
0.91	0.901	2	0.862	1	0.922	1	0.949	1	0.939	1
0.92	0.898	2	0.858	1	0.922	1	0.949	1	0.939	1
0.93	0.898	2	0.852	1	0.922	1	0.950	1	0.939	1
0.94	0.900	2	0.846	3	0.910	2	0.949	1	0.939	1
0.95	0.900	2	0.833	6	0.911	3	0.950	1	0.938	2
0.96	0.810	7	0.833	11	0.897	4	0.950	2	0.900	5
0.97	0.739	15	0.799	24	0.850	10	0.935	4	0.791	14
0.98	0.528	52	0.488	98	0.569	53	0.660	36	0.444	89
0.99	0.128	334	0.096	471	0.146	310	0.169	257	0.104	401

Table 7: All tested values for the threshold used for the Chinese Whispers clustering algorithm using the CFI dataset where ARI = 'adjusted rand score' and rel = the relative number of clusters.

params threshold	ArcFace		Dlib		FaceNet		face_recognition		VGG_Face	
	ARI	rel	ARI	rel	ARI	rel	ARI	rel	ARI	rel
0.70	0.000	-16	0.184	-15	0.153	-14	0.192	-15	0.199	-15
0.71	0.000	-16	0.184	-15	0.153	-14	0.192	-15	0.199	-15
0.72	0.000	-16	0.184	-15	0.153	-14	0.192	-15	0.199	-15
0.73	0.187	-15	0.234	-14	0.153	-14	0.192	-15	0.186	-15
0.74	0.000	-16	0.181	-15	0.265	-13	0.192	-15	0.199	-15
0.75	0.127	-15	0.234	-14	0.265	-13	0.192	-15	0.199	-15
0.76	0.217	-14	0.181	-15	0.265	-13	0.382	-14	0.376	-14
0.77	0.056	-14	0.232	-14	0.265	-13	0.355	-13	0.199	-15
0.78	0.121	-13	0.181	-15	0.430	-12	0.427	-13	0.386	-14
0.79	0.298	-12	0.232	-14	0.423	-12	0.427	-13	0.351	-14
0.80	0.162	-12	0.321	-14	0.276	-13	0.335	-14	0.386	-14
0.81	0.315	-11	0.331	-13	0.467	-11	0.484	-13	0.386	-14
0.82	0.475	-10	0.400	-13	0.449	-11	0.427	-13	0.345	-14
0.83	0.414	-10	0.374	-13	0.467	-11	0.619	-12	0.346	-14
0.84	0.419	-10	0.375	-12	0.460	-11	0.655	-11	0.516	-12
0.85	0.668	-8	0.640	-10	0.586	-10	0.745	-10	0.549	-11
0.86	0.551	-9	0.639	-10	0.668	-9	0.742	-10	0.736	-9
0.87	0.697	-7	0.617	-10	0.665	-9	0.745	-10	0.637	-10
0.88	0.693	-7	0.615	-10	0.729	-8	0.745	-10	0.743	-9
0.89	0.745	-6	0.618	-10	0.733	-8	0.775	-9	0.707	-10
0.90	0.744	-6	0.616	-10	0.707	-7	0.771	-9	0.764	-7
0.91	0.668	-3	0.707	-6	0.792	-6	0.856	-6	0.812	-4
0.92	0.679	-1	0.769	-3	0.767	-4	0.885	-4	0.812	-3
0.93	0.658	4	0.770	-2	0.752	-1	0.905	-3	0.773	-4
0.94	0.549	12	0.699	0	0.753	1	0.897	-3	0.731	2
0.95	0.319	26	0.483	15	0.726	5	0.797	1	0.707	9
0.96	0.219	48	0.424	35	0.421	23	0.721	8	0.429	24
0.97	0.126	100	0.271	61	0.267	59	0.531	24	0.265	57
0.98	0.045	199	0.082	160	0.101	137	0.243	71	0.096	131
0.99	0.014	288	0.014	285	0.023	260	0.044	205	0.022	256

Table 8: All tested values for the threshold used for the Chinese Whispers clustering algorithm using the Embedding on the Wall dataset where ARI = 'adjusted rand score' and rel = the relative number of clusters.